

Inference Latency/Throughput Scaling & Peak VRAM footprint

